

A SHORT INTRODUCTION TO PYTHON

Numpy, Scipy

Shai Fine

Slides adapted from a presentation by

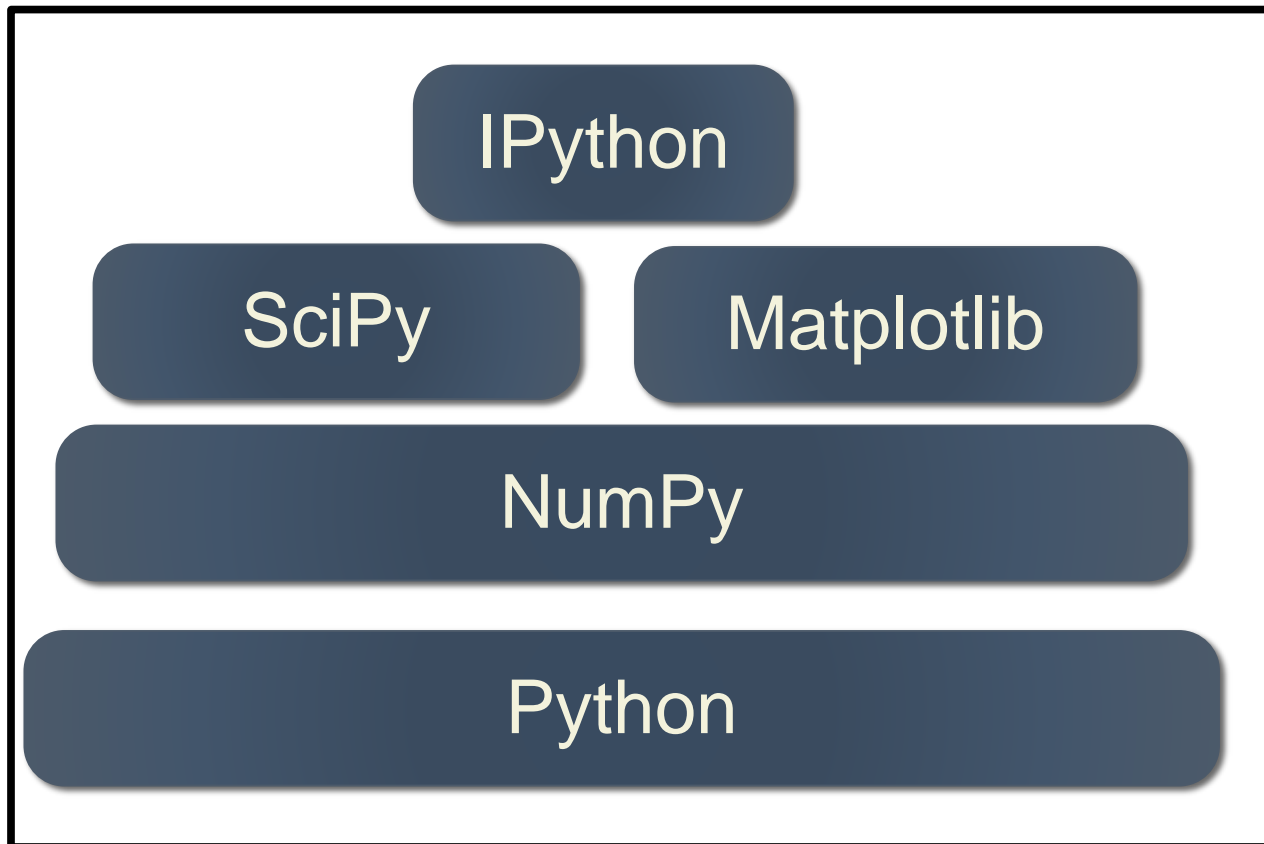


Essential Python Extensions

- The following packages extend Python with extra features
 - [NumPy](#) – Fast, multidimensional arrays
 - [SciPy](#) – Libraries of reliable, tested scientific functions
- Additional packages for Data Science (not covered today)
 - Wide range of learning algorithms ([scikit-learn](#))
 - Tools for data manipulation ([Pandas](#))
 - Plotting tools ([Matplotlib](#))
 - Direct connection to R ([rpy2](#))

PyLab

Sometimes the union of the 5 packages is called pylab



Helpful Sites

SCIPY DOCUMENTATION PAGE

<http://www.scipy.org/Documentation>

The screenshot shows the SciPy.org website. The header includes the SciPy.org logo and the text "Sponsored By ENTHOUGHT". The main content area is titled "Documentation" and includes a note about the "Installing SciPy" and "Cookbook" areas. Below this is a section titled "Getting Started and Tutorial" with a list of links: FAQ, Guide to NumPy, Numpy Glossary, Tentative NumPy Tutorial, Numpy Example List, Numpy Example List With Doc, Extensive NumPy & SciPy Summary, NumPy for MATLAB@ Users, RecordArrays, and Porting to NumPy. To the left of the main content is a sidebar with a "Wiki" section containing links to SciPy, Documentation, Mailing Lists, Download, Installing SciPy, Topical Software, Cookbook, Developer Zone, RecentChanges, and FindPage. Below the Wiki section is a "Page" section with links to Immutable Page, Info, Attachments, and a "More Actions" dropdown menu.

NUMPY EXAMPLES

http://www.scipy.org/Numpy_Example_List_With_Doc

The screenshot shows the "Numpy Example List With Doc" page. It features a sidebar with a "Wiki" section containing links to SciPy, Documentation, Mailing Lists, Download, Installing SciPy, Topical Software, Cookbook, Developer Zone, RecentChanges, and FindPage. The main content area is titled "Numpy Example List With Doc" and includes a note about the auto-generated version of the Numpy Example List. Below this is a section titled "Contents" with a list of links: 1. ..., 2. [], 3. T, 4. abs(), 5. absolute(), 6. accumulate, and 7. add().

[apply_along_axis\(\)](#)

`numpy.apply_along_axis(func1d, axis, arr, *args)`

Execute `func1d(arr[i],*args)` where `func1d` takes 1-D arrays and `arr` is an N-d array. `i` varies so as to apply the function along the given axis for each 1-d subarray in `arr`.

Example:

```
>>> from numpy import *
>>> def myfunc(a):
...     return (a[0]+a[-1])/2
...
>>> b = array([[1,2,3],[4,5,6],[7,8,9]])
>>> apply_along_axis(myfunc,0,b)
array([4.  5.  6])
>>> apply_along_axis(myfunc,1,b)
array([2.  5.  8])
```

NUMPY

Numerical Python

What is NumPy?

- NumPy is the fundamental package for scientific computing with Python
- NumPy provides a fast built-in object, `ndarray`, which is a **multi-dimensional array of a homogeneous data-type** that can be **manipulated in a vectorized form**
 - Numpy Offers Matlab-ish capabilities within Python
- NumPy can also be used as an efficient multi-dimensional container of generic data
 - This allows NumPy to seamlessly and speedily integrate with a wide variety of databases
- Website – <http://www.numpy.org/>
- Chronology
 - Initially developed by Travis Oliphant
 - NumPy 1.0 released October, 2006
 - ~20K downloads/month from Sourceforge
 - Doesn't count distributions that include NumPy
 - NumPy is at the core of nearly every scientific Python

Overview of NumPy

N-Dimensional ARRAY (NDARRAY)

- A NumPy array is a homogeneous collection of “items” of the same “data-type” (dtype)
 - Can be 1-dim or N-dims
- Element of the array can be C-structure or simple data-type
- Fast algorithms on machine data-types (int, float, etc.)

Universal Functions (UFUNC)

- Functions that operate element-by-element and return result
- Fast-loops registered for each fundamental data-type
 - $\sin(\mathbf{x}) = [\sin(x_i), i = 0 \dots N]$
 - $\mathbf{x} + \mathbf{y} = [x_i + y_i, i = 0 \dots N]$

Arrays in Python

- Python doesn't include a built-in multi-dimensional array
- Lists ok for storing small amounts of one-dimensional data

```
>>> a = [1,3,5,7,9]
>>> print(a[2:4])
[5, 7]
>>> b = [[1, 3, 5, 7, 9], [2, 4, 6, 8, 10]]
>>> print(b[0])
[1, 3, 5, 7, 9]
>>> print(b[1][2:4])
[6, 8]
```

```
>>> a = [1,3,5,7,9]
>>> b = [3,5,6,7,9]
>>> c = a + b
>>> print c
[1, 3, 5, 7, 9, 3, 5, 6, 7, 9]
```

- But, can't use directly with arithmetical operators (+, -, *, /, ...)
- Need efficient arrays with arithmetic and better multidimensional tools

Introducing NumPy Arrays

SIMPLE ARRAY CREATION

```
>>> a = array([0,1,2,3])
>>> a
array([0, 1, 2, 3])
```

CHECKING THE TYPE

```
>>> type(a)
<type 'array'>
```

NUMERIC 'TYPE' OF ELEMENTS

```
>>> a.dtype
dtype('int32')
```

NUMBER OF DIMENSIONS

```
>>> a.ndim
1
```

ARRAY SHAPE

```
# shape returns a tuple
# listing the length of the
# array along each dimension.
>>> a.shape
(4,)
>>> shape(a)
(4,)
```

ARRAY SIZE

```
# size reports the entire
# number of elements in an
# array.
>>> a.size
4
>>> size(a)
4
```

Introducing NumPy Arrays

ARRAY COPY

```
# create a copy of the array
>>> b = a.copy()
>>> b
array([0, 1, 2, 3])
```

CONVERSION TO LIST

```
# convert a numpy array to a
# python list
>>> a.tolist()
[0, 1, 2, 3]

# For 1D arrays, list also
# works equivalently, but
# is slower
>>> list(a)
[0, 1, 2, 3]
```

Setting Array Elements

ARRAY INDEXING

```
>>> a[0]
0
>>> a[0] = 10
>>> a
[10, 1, 2, 3]
```

FILL

```
# set all values in an array
>>> a.fill(0)
>>> a
[0, 0, 0, 0]

# This also works, but may
# be slower
>>> a[:] = 1
>>> a
[1, 1, 1, 1]
```



BEWARE OF TYPE COERSION

```
>>> a.dtype
dtype('int32')
```

```
# assigning a float to
# an int32 array will
# truncate decimal part
```

```
>>> a[0] = 10.6
>>> a
[10, 1, 2, 3]
```

```
# fill has the same behavior
```

```
>>> a.fill(-4.8)
>>> a
[-4, -4, -4, -4]
```

Multi-Dimensional Arrays (ndarray)

MULTI-DIMENSIONAL ARRAYS

```
>>> a = array([[ 0, 1, 2, 3],
               [10,11,12,13]])
```

```
>>> a
array([[ 0, 1, 2, 3],
       [10,11,12,13]])
```

(ROWS,COLUMNS)

```
>>> a.shape
(2, 4)
>>> shape(a)
(2, 4)
```

ELEMENT COUNT


```
>>> a.size
8
>>> size(a)
8
```

NUMBER OF DIMENSIONS

```
>>> a.ndim
2
```

GET/SET ELEMENTS

```
>>> a[1,3]
13
```



```
>>> a[1,3] = -1
>>> a
array([[ 0, 1, 2, 3],
       [10,11,12,-1]])
```

ADDRESS FIRST ROW USING SINGLE INDEX

```
>>> a[1]
array([10, 11, 12, -1])
```

Array Slicing

SLICING WORKS MUCH LIKE
STANDARD PYTHON SLICING

```
>>> a[0,3:5]  
array([3, 4])
```

```
>>> a[4:,4:]  
array([[44, 45],  
       [54, 55]])
```

```
>>> a[:,2]  
array([2, 12, 22, 32, 42, 52])
```

STRIDES ARE ALSO POSSIBLE

```
>>> a[2::2,::2]  
array([[20, 22, 24],  
       [40, 42, 44]])
```

0	1	2	3	4	5
10	11	12	13	14	15
20	21	22	23	24	25
30	31	32	33	34	35
40	41	42	43	44	45
50	51	52	53	54	55

Slices Are References

Slices are references to memory in original array. Changing values in a slice also changes the original array.

```
>>> a = array((0,1,2,3,4))

# create a slice containing only the
# last element of a
>>> b = a[2:4]
>>> b[0] = 10

# changing b changed a!
>>> a
array([ 1,  2, 10,  3,  4])
```

Fancy Indexing in 2D

Indexing by position

```
>>> a[(0,1,2,3,4), (1,2,3,4,5)]  
array([ 1, 12, 23, 34, 45])
```

```
>>> a[3:,[0, 2, 5]]  
array([[30, 32, 35],  
       [40, 42, 45]],  
       [50, 52, 55]])
```

Indexing with Booleans

```
>>> mask = array([1,0,1,0,0,1],  
                  dtype=bool)
```

```
>>> a[mask,2]  
array([2, 22, 52])
```

0	1	2	3	4	5
10	11	12	13	14	15
20	21	22	23	24	25
30	31	32	33	34	35
40	41	42	43	44	45
50	51	52	53	54	55



Unlike slicing, fancy indexing creates copies instead of views into original arrays.

Array Calculation Methods

SUM FUNCTION

```
>>> a = array([[1,2,3],  
               [4,5,6]], float)
```

```
# Sum defaults to summing all  
# *all* array values.
```

```
>>> sum(a)  
21.
```

```
# supply the keyword axis to  
# sum along the 0th axis.
```

```
>>> sum(a, axis=0)  
array([5., 7., 9.])
```

```
# supply the keyword axis to  
# sum along the last axis.
```

```
>>> sum(a, axis=-1)  
array([6., 15.])
```

SUM ARRAY METHOD

```
# The a.sum() defaults to  
# summing *all* array values
```

```
>>> a.sum()  
21.
```

```
# Supply an axis argument to  
# sum along a specific axis.
```

```
>>> a.sum(axis=0)  
array([5., 7., 9.])
```

PRODUCT

```
# product along columns.
```

```
>>> a.prod(axis=0)  
array([ 4., 10., 18.])
```

```
# functional form.
```

```
>>> prod(a, axis=0)  
array([ 4., 10., 18.])
```


Min/Max

MIN

```
>>> a = array([2.,3.,0.,1.])
>>> a.min(axis=0)
0.
# use Numpy's amin() instead
# of Python's builtin min()
# for speed operations on
# multi-dimensional arrays.
>>> amin(a, axis=0)
0.
```

ARGMIN

```
# Find index of minimum value.
>>> a.argmin(axis=0)
2
# functional form
>>> argmin(a, axis=0)
2
```

MAX

```
>>> a = array([2.,1.,0.,3.])
>>> a.max(axis=0)
3.
# functional form
>>> amax(a, axis=0)
3.
```

ARGMAX

```
# Find index of maximum value.
>>> a.argmax(axis=0)
1
# functional form
>>> argmax(a, axis=0)
1
```

Statistics Array Methods

MEAN

```
>>> a = array([[1,2,3],  
               [4,5,6]], float)
```

```
# mean value of each column
```

```
>>> a.mean(axis=0)  
array([ 2.5,  3.5,  4.5])
```

```
>>> mean(a, axis=0)  
array([ 2.5,  3.5,  4.5])
```

```
>>> average(a, axis=0)  
array([ 2.5,  3.5,  4.5])
```

```
# average can also calculate
```

```
# a weighted average
```

```
>>> average(a, weights=[1,2],  
            ...         axis=0)  
array([ 3.,  4.,  5.])
```

STANDARD DEV./VARIANCE

```
# Standard Deviation
```

```
>>> a.std(axis=0)  
array([ 1.5,  1.5,  1.5])
```

```
# Variance
```

```
>>> a.var(axis=0)  
array([2.25, 2.25, 2.25])
```

```
>>> var(a, axis=0)  
array([2.25, 2.25, 2.25])
```

Other Array Methods

CLIP

```
# Limit values to a range
>>> a = array([[1,2,3],
               [4,5,6]], float)

# Set values < 3 equal to 3.
# Set values > 5 equal to 5.
>>> a.clip(3,5)
array([[ 3.,  3.,  3.],
       [ 4.,  5.,  5.]])
```

POINT TO POINT

```
# Calculate max - min for
# array along columns
>>> a.ptp(axis=0)
array([ 3.0,  3.0,  3.0])
# max - min for entire array.
>>> a.ptp(axis=None)
5.0
```

ROUND

```
# Round values in an array.
# Numpy rounds to even, so
# 1.5 and 2.5 both round to 2.
>>> a = array([1.35, 2.5, 1.5])
>>> a.round()
array([ 1.,  2.,  2.])

# Round to first decimal place.
>>> a.round(decimals=1)
array([ 1.4,  2.5,  1.5])
```

Universal Functions (ufunc)

- ufuncs are objects that rapidly evaluate a function element-by-element over an array.
- Core piece is a 1-d loop written in C that performs the operation over the largest dimension of the array
- For 1-d arrays it is equivalent to but much faster than list comprehension

```
>>> type(np.exp)
<type 'numpy.ufunc'>
>>> x = array([1,2,3,4,5])
>>> print np.exp(x)
[2.71828, 7.38905, 20.08553, 54.59815, 148.41315]
>>> print [math.exp(val) for val in x]
[2.71828, 7.38905, 20.08553, 54.59815, 148.41315]
```

note: values reformatted to fit slide

Vectorizing Functions

VECTORIZING FUNCTIONS

Example


```
# special.sinc already available
# This is just for show.
def sinc(x):
    if x == 0.0:
        return 1.0
    else:
        w = pi*x
        return sin(w) / w
```

SOLUTION

```
>>> from numpy import vectorize
>>> vsinc = vectorize(sinc)
>>> vsinc([1.3,1.5])
array([-0.1981, -0.2122])
```

attempt

```
>>> sinc([1.3,1.5])
TypeError: can't multiply
sequence to non-int
>>> x = r_[-5:5:100j]
>>> y = vsinc(x)
>>> plot(x, y)
```



Mathematic Binary Operators *element by element*

$a + b \rightarrow \text{add}(a,b)$
 $a - b \rightarrow \text{subtract}(a,b)$
 $a \% b \rightarrow \text{remainder}(a,b)$

$a * b \rightarrow \text{multiply}(a,b)$
 $a / b \rightarrow \text{divide}(a,b)$
 $a ** b \rightarrow \text{power}(a,b)$

MULTIPLY BY A SCALAR

```
>>> a = array((1,2))
>>> a*3.
array([3., 6.] )
```

ELEMENT BY ELEMENT ADDITION

```
>>> a = array([1,2])
>>> b = array([3,4])
>>> a + b
array([4, 6])
```

ADDITION USING AN OPERATOR FUNCTION

```
>>> add(a,b)
array([4, 6])
```



IN PLACE OPERATION

```
# Overwrite contents of a.
# Saves array creation
# overhead
>>> add(a,b,a) # a += b
array([4, 6])
>>> a
array([4, 6])
```

Comparison and Logical Operators

<code>equal</code>	<code>(==)</code>	<code>not_equal</code>	<code>(!=)</code>	<code>greater</code>	<code>(>)</code>
<code>greater_equal</code>	<code>(>=)</code>	<code>less</code>	<code>(<)</code>	<code>less_equal</code>	<code>(<=)</code>
<code>logical_and</code>		<code>logical_or</code>		<code>logical_xor</code>	
<code>logical_not</code>					

2D EXAMPLE

```
>>> a = array((1,2,3,4), (2,3,4,5))
>>> b = array((1,2,5,4), (1,3,4,5))
>>> a == b
array([[True, True, False, True],
       [False, True, True, True]])
# functional equivalent
>>> equal(a,b)
array([[True, True, False, True],
       [False, True, True, True]])
```

Bitwise Operators *work only on Integer arrays*

<code>bitwise_and</code>	<code>(&)</code>	<code>invert</code>	<code>(~)</code>	<code>right_shift(a, shifts)</code>
<code>bitwise_or</code>	<code>()</code>	<code>bitwise_xor</code>		<code>left_shift(a, shifts)</code>

BITWISE EXAMPLES

```
>>> a = array((1,2,4,8))
>>> b = array((16,32,64,128))
>>> bitwise_or(a,b)
array([ 17,  34,  68, 136])
```

bit inversion

```
>>> a = array((1,2,3,4), uint8)
>>> invert(a)
array([254, 253, 252, 251], dtype=uint8)
```

left shift operation

```
>>> left_shift(a,3)
array([ 8, 16, 24, 32], dtype=uint8)
```


Matrix


- For two dimensional arrays NumPy defined a special matrix class in module `matrix`
 - Objects are created either with `matrix()` or `mat()` or converted from an array with method `asmatrix()`

```
>>> import numpy
>>> m = numpy.mat([[1,2],[3,4]])
# or
>>> a = numpy.array([[1,2],[3,4]])
>>> m = numpy.mat(a)
# or
>>> a = numpy.array([[1,2],[3,4]])
>>> m = numpy.asmatrix(a)
```


- Note that the statement `m = mat(a)` creates a copy of array 'a', whereas, method `m = asmatrix(a)` returns a new reference to the same data

Broadcasting

- Multiple inputs must be “broadcasted” to the same shape
 - All arrays are promoted to the same number of dimensions
 - All dimensions of length 1 are expanded as needed



- The *trailing* axes of both arrays must either be 1 or have the same size for broadcasting to occur



Matrix Objects

STRING CONSTRUCTION

```
>>> from numpy import mat
>>> a = mat('[1,3,5;2,5,1;2,3,6]')
>>> a
matrix([[1, 3, 5],
        [2, 5, 1],
        [2, 3, 6]])
```

TRANPOSE ATTRIBUTE

```
>>> a.T
matrix([[1, 2, 2],
        [3, 5, 3],
        [5, 1, 6]])
```

INVERTED ATTRIBUTE

```
>>> a.I
matrix([[ -1.1739,  0.1304,  0.956],
        [ 0.4347,  0.1739, -0.391],
        [ 0.1739, -0.130,  0.0434]
        ])
```

DIAGONAL

```
>>> a.diagonal()
matrix([[1, 5, 6]])
>>> a.diagonal(-1)
matrix([[3, 1]])
```

SOLVE

```
>>> b = mat('10;8;3')
>>> a.I*b
matrix([[ -7.82608696],
        [ 4.56521739],
        [ 0.82608696]])
```

```
>>> from scipy import linalg
>>> linalg.solve(a,b)
matrix([[ -7.82608696],
        [ 4.56521739],
        [ 0.82608696]])
```

Matrix vs. Array

- Operator `*`, `dot()`, and `multiply()`:
 - Array – `'*` means element-wise multiplication; `dot()` is used for matrix mul.
 - Matrix – `'*` means matrix multiplication; `multiply()` is used for element-wise mul.
- Handling of vectors (rank-1 arrays)
 - Array – the vector shapes `1xN`, `Nx1` are different things. Operations like `A[:,1]` return a rank-1 of shape `N`, not a rank-2 of shape `Nx1`. Transpose a rank-1 array does nothing
 - Matrix – rank-1 arrays are always upgraded to `1xN` or `Nx1` matrices (row or column vectors). `A[:,1]` returns a rank-2 matrix of shape `Nx1`
- Handling of higher-rank arrays (rank > 2)
 - Array objects can have rank > 2
 - Matrix objects always have exactly rank 2
- Convenience attributes
 - Array has a `.T` attribute, which returns the transpose of the data
 - Matrix has `.T`, `.H`, `.I`, and `.A` attributes, which return the conjugate transpose, inverse, and `asarray()` of the matrix, respectively.
- Convenience constructor
 - Array constructor takes (nested) Python sequences as initializers
 - Matrix constructor additionally takes a convenient string initializer

Example – Array and Matrix Calc.

```
>>> A = np.array([[n+m*10 for n in range(5)] for m in range(5)])
>>> v1 = arange(0, 5)
>>> A
array([[ 0,  1,  2,  3,  4],
       [10, 11, 12, 13, 14],
       [20, 21, 22, 23, 24],
       [30, 31, 32, 33, 34],
       [40, 41, 42, 43, 44]])
>>> v1
array([0, 1, 2, 3, 4])
>>> np.dot(A,A)
array([[ 300,  310,  320,  330,  340],
       [1300, 1360, 1420, 1480, 1540],
       [2300, 2410, 2520, 2630, 2740],
       [3300, 3460, 3620, 3780, 3940],
       [4300, 4510, 4720, 4930, 5140]])
>>> np.dot(A,v1)
array([ 30, 130, 230, 330, 430])
>>> np.dot(v1,v1)
30
```

Examples – Array and Matrix Calc.

```
# Alternatively, we can cast the array objects to the type
# matrix. This # changes the behavior of the standard
# arithmetic operators +, -, * to # use matrix algebra.
>>> M = np.matrix(A)
>>> v = np.matrix(v1).T
>>> v
matrix([[0],
        [1],
        [2],
        [3],
        [4]])
>>> M*v
matrix([[ 30],
        [130],
        [230],
        [330],
        [430]])
>>> v.T * v    # inner product
matrix([[30]])
```

Concluding Remarks

- Using arrays wisely
 - Array operations are implemented in C or Fortran
 - Optimized algorithms - i.e. fast!
 - Python loops (i.e. for i in a:...) are much slower
 - Prefer array operations over loops, especially when speed important
 - Also produces shorter code, often more readable
- Matrix or Array, which one to use?
 - Short answer – Use Array
 - They are the standard vector/matrix/tensor type of NumPy. Many NumPy functions return arrays, not matrices
 - There is a clear distinction between element-wise and linear algebra operations
 - You can have standard vectors or row/column vectors if you like
 - The main disadvantage of using the array type is that you will have to use `dot()` instead of `*` matrix multiplication
- NumPy for Matlab Users
 - http://wiki.scipy.org/NumPy_for_Matlab_Users

SCIPY

Scientific Python

SciPy Overview

- Available at www.scipy.org


CURRENT PACKAGES

- Special Functions (scipy.special)
 - Signal Processing (scipy.signal)
 - Image Processing (scipy.ndimage)
 - Fourier Transforms (scipy.fftpack)
 - Optimization (scipy.optimize)
 - Numerical Integration (scipy.integrate)
 - Linear Algebra (scipy.linalg)
- Input/Output (scipy.io)
 - Statistics (scipy.stats)
 - Fast Execution (scipy.weave)
 - Clustering Algorithms (scipy.cluster)
 - Sparse Matrices (scipy.sparse)
 - Interpolation (scipy.interpolate)
 - More (e.g. scipy.odr, scipy.maxentropy)

Image Processing

```
# The famous lena image is packaged with scipy
>>> from scipy import lena, signal
>>> lena = lena().astype(float32)
>>> imshow(lena, cmap=cm.gray)
# Blurring using a median filter
>>> f1 = signal.medfilt2d(lena, [15,15])
>>> imshow(f1, cmap=cm.gray)
```

LENA IMAGE



MEDIAN FILTERED IMAGE






Image Processing

```
# Edge detection using Sobel filter  
>>> from scipy.ndimage.filters import sobel  
>>> imshow(lena)  
>>> edges = sobel(lena)  
>>> imshow(edges)
```

NOISY IMAGE



FILTERED IMAGE



Statistics

scipy.stats --- CONTINUOUS DISTRIBUTIONS

over 80
continuous
distributions!

METHODS










pdf

cdf

rvs

ppf

stats



Statistics

scipy.stats --- Discrete Distributions

10 standard
discrete
distributions
(plus any
arbitrary
finite RV)

METHODS




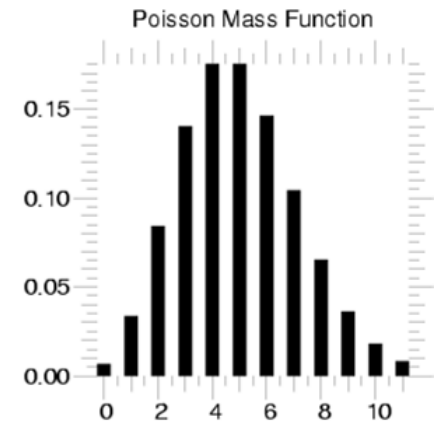

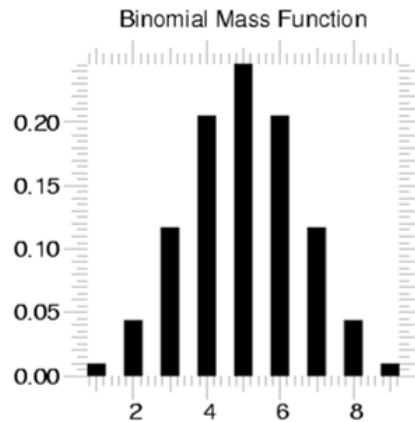
pdf

cdf

rvs

ppf

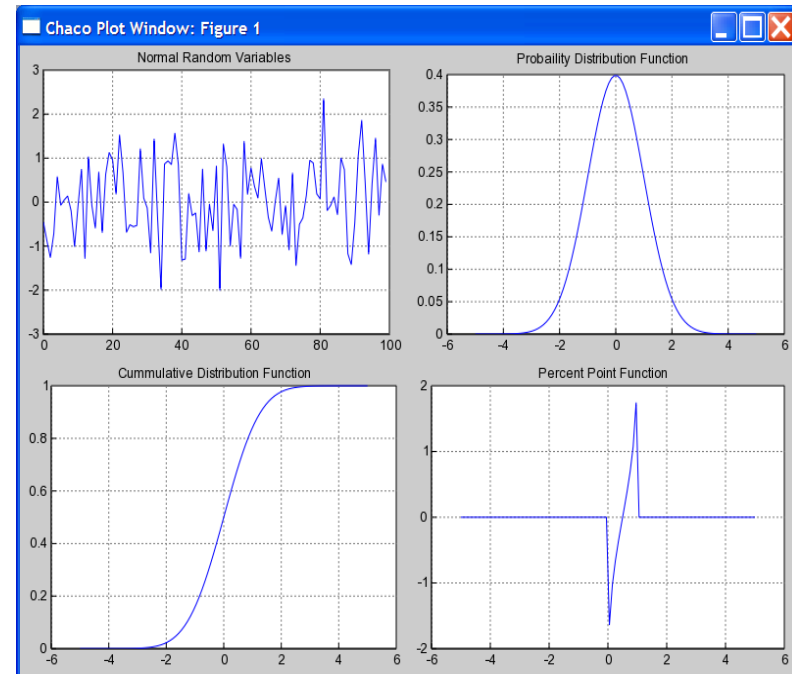
stats



Using Stats Objects

DISTRIBUTIONS

```
# Sample normal dist. 100 times.  
>>> samp = stats.norm.rvs(size=100)  
  
>>> x = r_[-5:5:100j]  
# Calculate probability dist.  
>>> pdf = stats.norm.pdf(x)  
# Calculate cumulative Dist.  
>>> cdf = stats.norm.cdf(x)  
# Calculate Percent Point Function  
>>> ppf = stats.norm.ppf(x)
```



Statistics

scipy.stats --- Basic Statistical Calculations on Data

- `numpy.mean`, `numpy.std`, `numpy.var`, `numpy.cov`
- `stats.skew`, `stats.kurtosis`, `stats.moment`

scipy.stats.bayes_mvs --- Bayesian mean, variance, and std.

```
# Create "frozen" Gamma distribution with a=2.5
>>> grv = stats.gamma(2.5)
>>> grv.stats()      # Theoretical mean and variance
(array(2.5), array(2.5))
# Estimate mean, variance, and std with 95% confidence
>>> vals = grv.rvs(size=100)
>>> stats.bayes_mvs(vals, alpha=0.95)
((2.52887906081, (2.19560839724, 2.86214972438)),
 (2.87924964268, (2.17476164549, 3.8070215789)),
 (1.69246760584, (1.47470730841, 1.95115903475)))
# (expected value and confidence interval for each of
# mean, variance, and standard-deviation)
```

Statistics

Continuous PDF Estimation using Gaussian Kernel Density Estimation

```
# Sample normal dist. 100 times.
```

```
>>> rv1 = stats.norm()
```

```
>>> rv2 = stats.norm(2.0,0.8)
```

```
>>> samp = r_[rv1.rvs(size=100),  
               rv2.rvs(size=100)]
```

```
# Kernel estimate (smoothed histogram)
```


```
>>> apdf = stats.kde.gaussian_kde(samp)
```

```
>>> x = linspace(-3,6,200)
```

```
>>> plot(x, apdf(x), 'r')
```

```
# Histogram
```

```
>>> hist(x, bins=25, normed=True)
```



Linear Algebra

scipy.linalg --- FAST LINEAR ALGEBRA

- Uses ATLAS if available --- very fast
- Low-level access to BLAS and LAPACK routines in modules `linalg.fblas`, and `linalg.flapack` (FORTRAN order)
- High level matrix routines
 - Linear Algebra Basics: `inv`, `solve`, `det`, `norm`, `lstsq`, `pinv`
 - Decompositions: `eig`, `lu`, `svd`, `orth`, `cholesky`, `qr`, `schur`
 - Matrix Functions: `expm`, `logm`, `sqrtm`, `cosm`, `coshm`, `funm` (general matrix functions)

Linear Algebra

LU FACTORIZATION

```
>>> from scipy import linalg
>>> a = array([[1,3,5],
...           [2,5,1],
...           [2,3,6]])
# time consuming factorization
>>> lu, piv = linalg.lu_factor(a)

# fast solve for 1 or more
# right hand sides.
>>> b = array([10,8,3])
>>> linalg.lu_solve((lu, piv), b)
array([-7.82608696,  4.56521739,
        0.82608696])
```

EIGEN VALUES AND VECTORS

```
>>> from scipy import linalg
>>> a = array([[1,3,5],
...           [2,5,1],
...           [2,3,6]])
# compute eigen values/vectors
>>> vals, vecs = linalg.eig(a)
# print eigen values
>>> vals
array([ 9.39895873+0.j,
       -0.73379338+0.j,
        3.33483465+0.j])
# eigen vectors are in columns
# print first eigen vector
>>> vecs[:,0]
array([-0.57028326,
       -0.41979215,
       -0.70608183])
# norm of vector should be 1.0
>>> linalg.norm(vecs[:,0])
1.0
```

Optimization

scipy.optimize --- unconstrained minimization and root finding

- **Unconstrained Optimization**

`fmin` (Nelder-Mead simplex), `fmin_powell` (Powell's method), `fmin_bfgs` (BFGS quasi-Newton method), `fmin_ncg` (Newton conjugate gradient), `leastsq` (Levenberg-Marquardt), `anneal` (simulated annealing global minimizer), `brute` (brute force global minimizer), `brent` (excellent 1-D minimizer), `golden`, `bracket`

- **Constrained Optimization**

`fmin_l_bfgs_b`, `fmin_tnc` (truncated newton code), `fmin_cobyla` (constrained optimization by linear approximation), `fminbound` (interval constrained 1-d minimizer)

- **Root finding**


`fsolve` (using MINPACK), `brentq`, `brenth`, `ridder`, `newton`, `bisect`, `fixed_point` (fixed point equation solver)

Optimization

EXAMPLE: Non-linear least-squares data fitting

```
# fit data-points to a curve
# demo/data_fitting/datafit.py

>>> from numpy.random import randn
>>> from numpy import exp, sin, pi
>>> from numpy import linspace
>>> from scipy.optimize import leastsq
>>> def func(x,A,a,f,phi):
    return A*exp(-a*sin(f*x+pi/4))
>>> def errfunc(params, x, data):
    return func(x, *params) - data
>>> ptrue = [3,2,1,pi/4]
>>> x = linspace(0,2*pi,25)
>>> true = func(x, *ptrue)
>>> noisy = true + 0.3*randn(len(x))
>>> p0 = [1,1,1,1]
>>> pmin, ier = leastsq(errfunc, p0, args=(x, noisy))
>>> pmin
array([3.1705, 1.9501, 1.0206, 0.7034])
```



THANK YOU
